

# CENG 484 Data Mining

# **Data Mining with Python and R**

---

Asst. Prof. Serap Şahin

Res. Asst. Altuğ Yiğit

March 2020

---


# Introduction

This document has been prepared for the CENG 484 data mining course. It contains the necessary steps for processing the data with **Python** and **R**. These two programming languages are explained by comparison. Initially, document shows how to download and install Python, PyCharm, R and RStudio. After the environment is ready, it covers how to read, plot and analyze data. The main purpose of this document is to provide the necessary infrastructure to solve real world problems using mathematical data mining methods.

Data mining is the computational technique that enables us to **find patterns** and learn some rules hidden in data sets. It is an interdisciplinary field with contributions from many areas, such as statistics, machine learning, information retrieval, pattern recognition and bioinformatics. Data mining is widely used in many domains, such as retail, finance, telecommunication and social media. The main techniques for data mining include **classification** and prediction, **clustering**, outlier detection, association rules, sequence analysis, time series analysis and text mining, and also some new techniques such as social network analysis and sentiment analysis.


In real world applications, a data mining process can be broken into **six major phases**: business understanding, data understanding, data preparation, modeling, evaluation and deployment, as defined by the **CRISP-DM** (Cross Industry Standard Process for Data Mining). These phases are shown in the figure below. One of the most important distinguishing issues in data mining is **size**. One has to consider issues like computational **efficiency**, limited **memory** resources, interfaces to databases, etc. All these issues turn data mining into a highly interdisciplinary subject involving tasks

**not only of typical data analysts** but also of people working with databases, data visualization on high dimensions, etc.



**Tools** such as RapidMiner, Orange, Weka, Knime and **programming languages** such as Python, R, Julia, Java are widely used for the application of data mining methods. This course focuses on Python and R, the most used languages.

**Python** is a high-level, interpreted, interactive and object-oriented scripting language. Python is designed to be highly readable. Python was developed by Guido van Rossum in the late 1980s and early 1990s at the National Research Institute for Mathematics and Computer Science in the



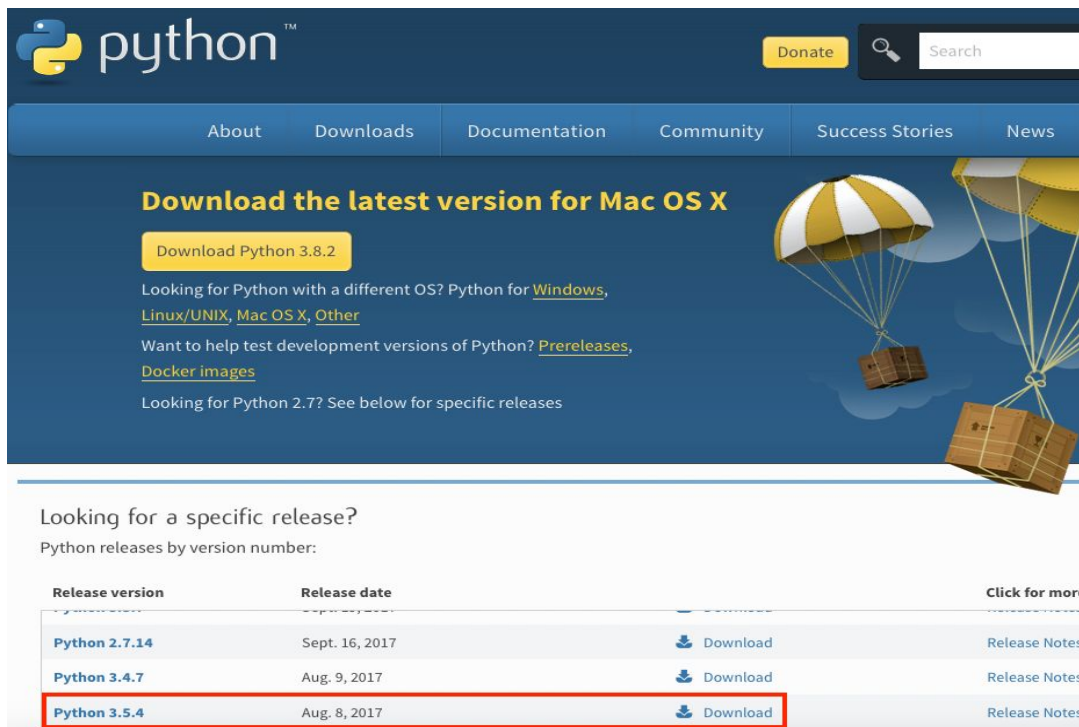
Netherlands. Python is derived from many other languages, including ABC, Modula-3, C, C++, Algol-68, SmallTalk, and Unix shell and other scripting languages. Python is copyrighted. Like Perl, Python source code is now available **under the General Public License** (GPL). Python 1.0 was released in November 1994.

**R** is a programming language and an environment for statistical computing. It is based on the computer language S, developed by John Chambers and others at Bell Laboratories in 1976. R was initially developed in 1996 by Ihaka and Gentleman, both from the University of Auckland, New Zealand. The source code of every R component is freely available for inspection and adaptation. This fact allows you to check and test the reliability of anything you use in R. Many classical and modern **statistical techniques** have been implemented in the R environment. A few of these are built into the **base** R environment, but many are supplied as **packages**.

## Part 1 - Preparing the Environment

### 1.1 Installing Python

Python is available for your system from the website (<https://www.python.org/downloads/>). In this course, **Python 3.5.4** version will be used, so this version may be selected from the website before downloading.



python™

Donate Search

About Downloads Documentation Community Success Stories News

### Download the latest version for Mac OS X

[Download Python 3.8.2](#)

Looking for Python with a different OS? Python for [Windows](#), [Linux/UNIX](#), [Mac OS X](#), [Other](#)

Want to help test development versions of Python? [Prereleases](#), [Docker images](#)

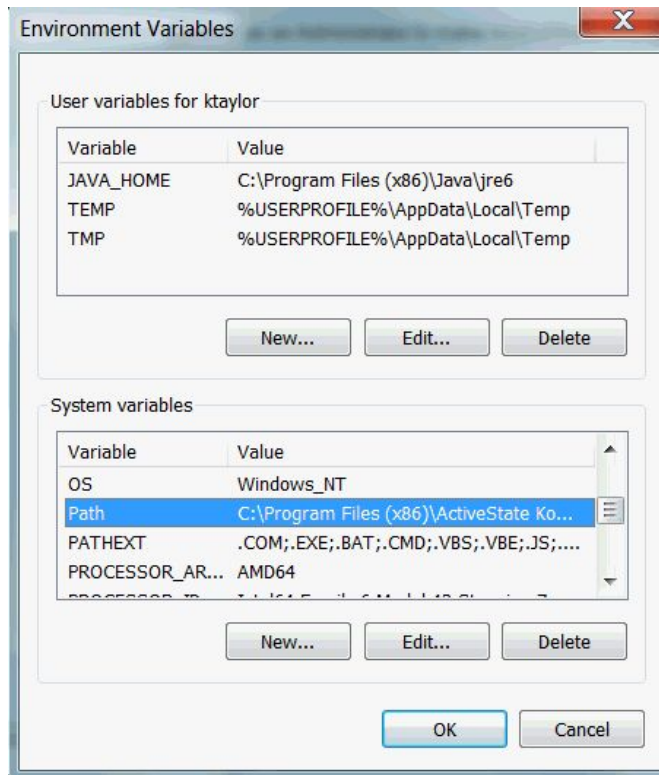
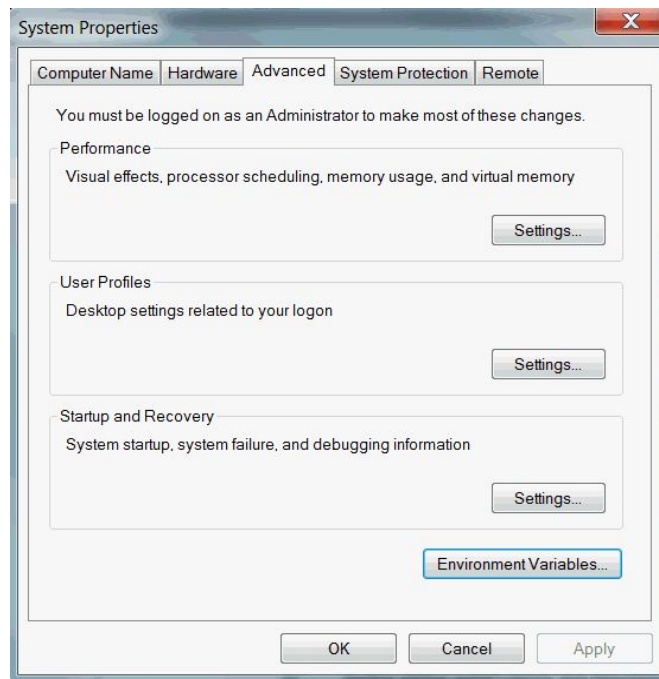
Looking for Python 2.7? See below for specific releases

Looking for a specific release?  
Python releases by version number:

| Release version               | Release date   |                          | Click for more                |
|-------------------------------|----------------|--------------------------|-------------------------------|
| <a href="#">Python 2.7.14</a> | Sept. 16, 2017 | <a href="#">Download</a> | <a href="#">Release Notes</a> |
| <a href="#">Python 3.4.7</a>  | Aug. 9, 2017   | <a href="#">Download</a> | <a href="#">Release Notes</a> |
| <a href="#">Python 3.5.4</a>  | Aug. 8, 2017   | <a href="#">Download</a> | <a href="#">Release Notes</a> |

If the installation is successful, when the command "**python --version**" is written from the console for **Unix** based systems, it gives information about the Python version. **Windows** users will need to set an environment variable to use **Python from the command line** as figures below. Find where you install Python onto your computer; the default location is "C:\Python35".

- Right-click the Computer icon and choose Properties,
- Choose Advanced system settings,
- On the Advanced tab, click Environment Variables,
- Click New to create a new environment variable. Click Edit to modify an existing environment variable,
- Add "C:\Python35" and "C:\Python35\Scripts" to path.



```
Command Prompt - python
Microsoft Windows [Version 10.0.18362.657]
(c) 2019 Microsoft Corporation. All rights reserved.

C:\Users\TR-OSF>python --version
Python 3.5.4

C:\Users\TR-OSF>python
Python 3.5.4 (v3.5.4:3f56838, Aug  8 2017, 02:17:05) [MSC v.1900 64 bit (AMD64)] on
win32
Type "help", "copyright", "credits" or "license" for more information.
>>> for i in range(0, 100): print(i)
...
_
```

## 1.2 Installing PyCharm


**PyCharm** is a cross-platform IDE that provides consistent experience on the Windows, macOS, and Linux operating systems. PyCharm is available in three editions: Professional, Community, and Edu. The Community and Edu editions are open-source projects and they are free. In this course, the **PyCharm Community** version will be used since it is sufficient. It can be downloaded from website (<https://www.jetbrains.com/pycharm/download/>).

20  
years

Tools Languages Solutions Support Company

PyCharm

Coming in 2020.1 What's New Features Learning Center Buy



Version: 2019.3.3  
Build: 193.6494.30  
7 February 2020

[System requirements](#)  
[Installation Instructions](#)  
[Other versions](#)

### Download PyCharm

Windows **Mac** Linux

**Professional**  
For both Scientific and Web  
Python development. With HTML,  
JS, and SQL support.  

Download

Free trial

**Community**  
For pure Python development  

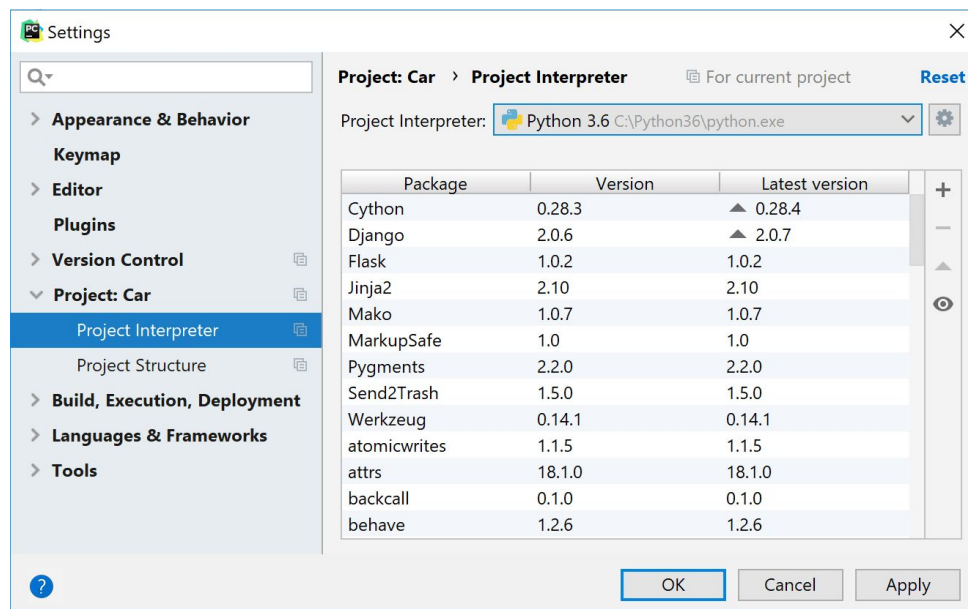
Download

Free, open-source

6



If you're on the Welcome screen, click Create New Project. In this screen you may create virtual environment, if you would like to **use environment** that is **on your computer**, you should select "**Existing interpreter**". To manage **Python packages** for the project interpreter in **PyCharm**, select the **Project Interpreter** page in the project **Settings/Preferences** or select Interpreter Settings in the Python Interpreter widget.



You can use **pip on console** to install packages from the Python Package Index and other indexes, pip is the package installer for Python. You may type package name that you want to install as "**pip install package\_name**".

```

(c) 2019 Microsoft Corporation. All rights reserved.

C:\Users\TR-OSF>pip install pandas
Collecting pandas
  Downloading https://files.pythonhosted.org/packages/7b/e4/ba4a6408fbc4f66bc7c84d15787fd3d50aa8ef11e47273e8a29d535d25ee/pandas-0.25.3-cp35-cp35m-win_amd64.whl (8.8MB)
    100% |#####| 8.8MB 133kB/s
Requirement already satisfied: python-dateutil>=2.6.1 in c:\python35\lib\site-packages (from pandas)
Requirement already satisfied: numpy>=1.13.3 in c:\python35\lib\site-packages (from pandas)
Collecting pytz>=2017.2 (from pandas)
  Downloading https://files.pythonhosted.org/packages/e7/f9/f0b53f88060247251bf481fa6ea62cd0d25bf1b11a87888e53ce5b7c8ad2/pytz-2019.3-py2.py3-none-any.whl (509kB)
    100% |#####| 512kB 936kB/s
Requirement already satisfied: six>=1.5 in c:\python35\lib\site-packages (from python-dateutil>=2.6.1->pandas)
Installing collected packages: pytz, pandas
Successfully installed pandas-0.25.3 pytz-2019.3
You are using pip version 9.0.1, however version 20.0.2 is available.
You should consider upgrading via the 'python -m pip install --upgrade pip' command.
  
```



## 1.3 Installing R

In order to install R in your system, the easiest way is to obtain a **binary distribution** from the R website (<https://cran.r-project.org>). This site is referred as **CRAN** (Comprehensive R Archive Network). Most users download and install a binary version from the site. Binary version is a different type of version that has been translated (by compilers) into machine language for execution on a given operating system. R is designed to be very **portable**: it will run on Microsoft Windows, Linux, Solaris, Mac OSX, and other operating systems, there are different binary versions for each. For this course, R 3.5.2 should be downloaded as shown in figures below.



CRAN

[Mirrors](#)

[What's new?](#)

[Task Views](#)

[Search](#)

About R

[R Homepage](#)

[The R Journal](#)

Software

[R Sources](#)

[R Binaries](#)

[Packages](#)

[Other](#)

Documentation

[Manuals](#)

[FAQs](#)

[Contributed](#)

### The Comprehensive R Archive Network

#### Download and Install R

Precompiled binary distributions of the base system and contributed packages, **Windows and Mac** users most likely want one of these versions of R:

- [Download R for Linux](#)
- [Download R for \(Mac\) OS X](#)
- [Download R for Windows](#)

R is part of many Linux distributions, you should check with your Linux package management system in addition to the link above.

#### Source Code for all Platforms

Windows and Mac users most likely want to download the precompiled binaries listed in the upper box, not the source code. The sources have to be compiled before you can use them. If you do not know what this means, you probably do not want to do it!

- The latest release (2020-02-29, Holding the Windsock) [R-3.6.3.tar.gz](#), read [what's new](#) in the latest version.
- Sources of [R alpha and beta releases](#) (daily snapshots, created only in time periods before a planned release).
- Daily snapshots of current patched and development versions are [available here](#). Please read about [new features and bug fixes](#) before filing corresponding feature requests or bug reports.
- Source code of older versions of R is [available here](#).
- Contributed extension [packages](#)

#### Questions About R

- If you have questions about R like how to download and install the software, or what the license terms are, please read our [answers to frequently asked questions](#) before you send an email.



CRAN  
[Mirrors](#)  
[What's new?](#)  
[Task Views](#)  
[Search](#)

*About R*  
[R Homepage](#)  
[The R Journal](#)

*Software*  
[R Sources](#)  
[R Binaries](#)  
[Packages](#)  
[Other](#)

*Documentation*  
[Manuals](#)  
[FAQs](#)  
[Contributed](#)

## R for Windows

### Subdirectories:

[base](#)

[contrib](#)

[old contrib](#)

[Rtools](#)

Binaries for base distribution. This is what you want to [install R for the first time](#).

Binaries of contributed CRAN packages (for R >= 2.13.x; managed by Uwe Ligges). There is also information on [third party software](#) available for CRAN Windows services and corresponding environment and make variables.

Binaries of contributed CRAN packages for outdated versions of R (for R < 2.13.x; managed by Uwe Ligges).

Tools to build R and R packages. This is what you want to build your own packages on Windows, or to build R itself.

Please do not submit binaries to CRAN. Package developers might want to contact Uwe Ligges directly in case of questions / suggestions related to Windows binaries.

You may also want to read the [R FAQ](#) and [R for Windows FAQ](#).

Note: CRAN does some checks on these binaries for viruses, but cannot give guarantees. Use the normal precautions with downloaded executables.

## R-3.6.3 for Windows (32/64 bit)

[Download R 3.6.3 for Windows](#) (83 megabytes, 32/64 bit)

[Installation and other instructions](#)

[New features in this version](#)

If you want to double-check that the package you have downloaded matches the package distributed by CRAN, you can compare the [md5sum](#) of the .exe to the [fingerprint](#) on the master server. You will need a version of md5sum for windows: both [graphical](#) and [command line versions](#) are available.

### Frequently asked questions

- [Does R run under my version of Windows?](#)
- [How do I update packages in my previous version of R?](#)
- [Should I run 32-bit or 64-bit R?](#)

Please see the [R FAQ](#) for general information about R and the [R Windows FAQ](#) for Windows-specific information.

### Other builds

- Patches to this release are incorporated in the [r-patched snapshot build](#).
- A build of the development version (which will eventually become the next major release of R) is available in the [r-devel snapshot build](#).
- [Previous releases](#)

Note to webmasters: A stable link which will redirect to the current Windows binary release is  
[<CRAN MIRROR>/bin/windows/base/release.htm](#)

Last change: 2020-02-29

## Previous Releases of R for Windows

This directory contains previous binary releases of R to run on Windows 95, 98, ME, NT4.0, 2000 and XP or later on Intel/clone chips.

The current release, and links to development snapshots, are available [here](#). Source code for these releases and others is available through [the main CRAN page](#).

In this directory:

[R 3.6.2](#) (December, 2019)

[R 3.6.1](#) (July, 2019)

[R 3.6.0](#) (April, 2019)

[R 3.5.3](#) (March, 2019)

[R 3.5.2](#) (December, 2018)

[R 3.5.1](#) (July, 2018)

[R 3.5.0](#) (April, 2018)

[R 3.4.4](#) (March, 2018)


[R 3.4.3](#) (November, 2017)

[R 3.4.2](#) (September, 2017)

[R 3.4.1](#) (June, 2017)

[R 3.4.0](#) (April, 2017)

[R 3.3.3](#) (March, 2017)



To run R in **Windows** you simply double-click the **R application** that is called like "R\_x64\_3.5.2.exe". In **Unix** versions you should **type R** at the operating **system prompt**. Both will bring up the R console with its prompt ">". Whenever you see this **R command prompt**, ">" you can interpret it as R waiting for you to enter a command. You type in the commands at the prompt and then press the enter key to ask R to execute them. This may or may not produce some form of output and then a new prompt appears. At the prompt you may use the arrow keys to browse and edit previously entered commands. This is handy when you want to type commands similar to what you have done before as you avoid typing them again.

If you want to **quit R** you can issue the command "**q()**" at the prompt. You will be asked if you want to save the current workspace. You should answer yes only if you want to resume your current analysis at the point you are leaving it, later on.

## 1.4 Installing RStudio

RStudio is a free and open source IDE (integrated development environment) for R. It is a very useful and powerful tool for R programming. **RStudio Desktop** open source edition can be downloaded from website as shown in figure below (<https://rstudio.com/products/rstudio/download/>).

## Choose Your Version

RStudio is a set of integrated tools designed to help you be more productive with R. It includes a console, syntax-highlighting editor that supports direct code execution, and a variety of robust tools for plotting, viewing history, debugging and managing your workspace.

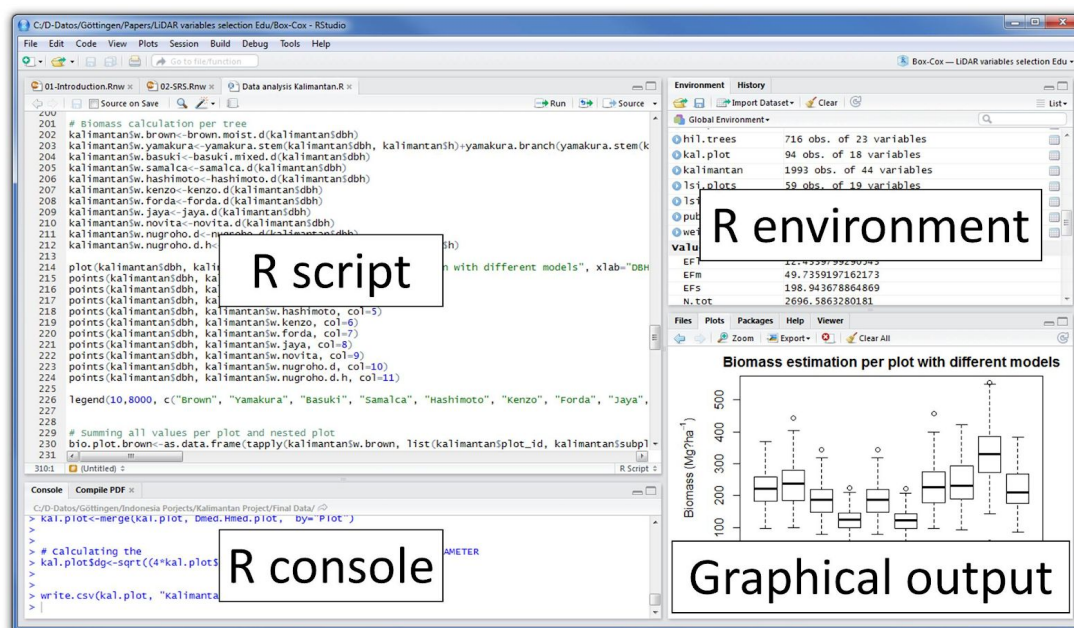
[LEARN MORE ABOUT RSTUDIO FEATURES](#)



|  |   |   |  |
|--|---|---|--|
| <b>RStudio Desktop</b><br>Open Source License<br><b>Free</b><br><a href="#">Download</a><br><a href="#">Learn more</a> | <b>RStudio Desktop</b><br>Commercial License<br><b>\$995 /year</b><br><a href="#">Buy</a><br><a href="#">Learn more</a> | <b>RStudio Server</b><br>Open Source License<br><b>Free</b><br><a href="#">Download</a><br><a href="#">Learn more</a> | <b>RStudio Server Pro</b><br>Commercial License<br><b>\$4,975 /year</b><br>(5 Named Users)<br><a href="#">Buy</a><br><a href="#">Evaluation</a>   <a href="#">Learn more</a> |
|--|---|---|--|

It is free of charge and can run on various operating systems like Windows, Mac and Linux. Your operating system will be automatically recognized by this website. Once the installation of R has completed successfully, run the RStudio installer. When RStudio is launched for the first time, you can see a similar window in figure below. There are **four** panels:

- **Source panel (top left)**, which shows your R source code.
- **Console panel (bottom left)**, which shows outputs and system messages displayed in a normal R console;
- **Environment/History/Presentation panel (top right)**, whose three tabs show respectively all objects and function loaded in R, a history of submitted R code, and Presentations generated with R;
- **Files/Plots/Packages/Help/Viewer panel (bottom right)**, whose tabs show respectively a list of files, plots, R packages installed, help documentation and local web content.



**Note:** If you cannot see the source panel, you can find it by clicking menu “File”, “New File” and then “R Script”. You can run a line or a selection of R code by clicking the “Run” button on top of source panel, or pressing “Ctrl + Enter”.

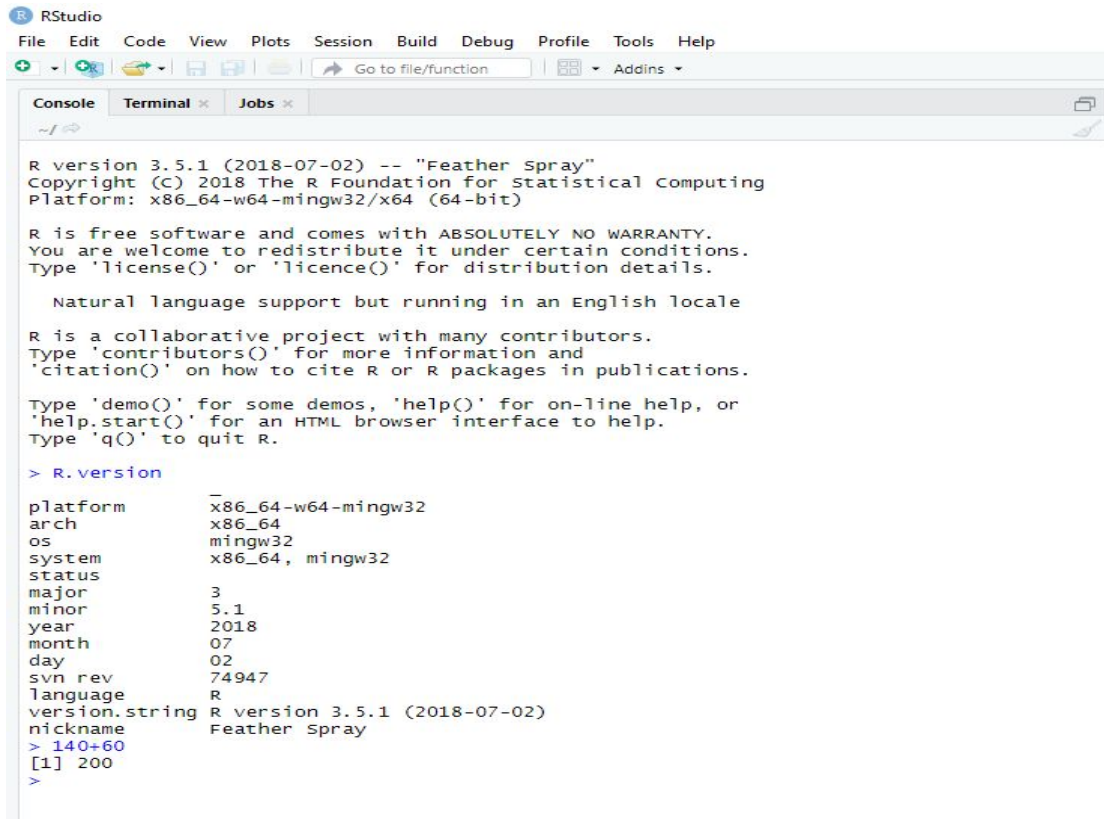
To create a **new project**, click the “Project” button at the top right corner and then choose “New Project”. After that, select “create project from new directory” and then “Empty Project”. After typing a directory name, which will also be your project name, click “Create Project” to create your project folder and files. After that, create three folders as:

- **code**, where to put your R source code
- **data**, where to put your datasets;
- **figures**, where to put produced diagrams.

"**R.version**" command may be written from the console screen to check the **R version via RStudio**. Some **arithmetic** operations such as “140+60” can be done on the console to check whether R Studio is properly working or not



as shown in figure below. The [1] that prefixes the output indicates that this is item 1 in a vector of output.



```
RStudio
File Edit Code View Plots Session Build Debug Profile Tools Help
Go to file/function Addins

Console Terminal Jobs

R version 3.5.1 (2018-07-02) -- "Feather Spray"
Copyright (C) 2018 The R Foundation for Statistical Computing
Platform: x86_64-w64-mingw32/x64 (64-bit)

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

Natural language support but running in an English locale

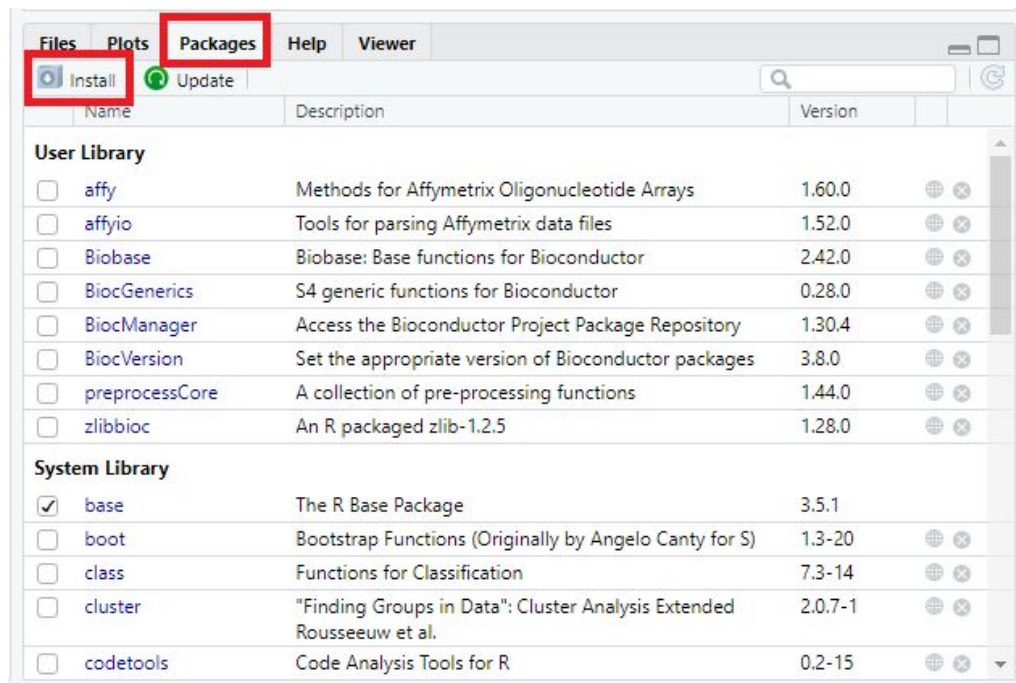
R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

> R.version
platform      x86_64-w64-mingw32
arch           x86_64
os            mingw32
system        x86_64, mingw32
status
major          3
minor          5.1
year          2018
month          07
day            02
svn rev       74947
language       R
version.string R version 3.5.1 (2018-07-02)
nickname       Feather Spray
> 140+60
[1] 200
>
```

To install a **R package**, come to the "**Packages**" tab and select "**Install**". Suppose we downloaded the `ggplot2` package, "**library (ggplot2)**" command may be written to check that this library is installed. If there is no message like "**Error in library(x) : there is no package called x**", it was successfully installed.

As a second option, the "**install.packages("ggplot2")**" command may be written manually on the console to install packages.



## Part 2 - Data Types

### 2.1 Python

Python built in main data types are: Text (String), Numeric (Integer, Float, Long), Complex, Boolean (True / False).

The most common data structure in Python is known as **list**.

```
a = [1.8, 4.5, 6.5] #float
b = [1 + 2*i, 3 - 6*i] #complex
d = [23, 44, 65] #integer
e = [True, False, True, False, True] #boolean
```



Python also includes different data structures, which are **dictionary** and **tuple**. Python **dictionary** is a collection of **key and value pairs** separated by a colon (:).

```
dict = {'no': '20151231212', 'name': 'Ali', 'age': 20 } # dictionary

print("Student no:", dict['no'])

print("Student name:", dict['name'])

print("Student age:", dict['age'])
```

Iterating over the elements of a **tuple is faster** compared to iterating over a **list**. We can **have tuple** of same type of data items as well as **mixed type** of data items.

```
d = (23, 44, 65) # tuple

e = (23, 60.4, "Ali") # mixed tuple

f = ("Ali", [80, 85, 90]) # mixed tuple
```

## 2.2 R

Everything you see or create in R is an **object**. R has 5 basic classes of objects: Character, Numeric (Real or Decimal Numbers), Integer (Whole Numbers), Complex, Logical (True / False). The most common data structure in R is known as **vector**. You can create an empty vector using `vector()`. A vector contains object of same class. You can also create vector using `c()` or concatenate command.

```
> a <- c(1.8, 4.5) #numeric
> b <- c(1 + 2i, 3 - 6i) #complex
> d <- c(23, 44) #integer
> e <- vector("logical", length = 5)
```

R contains **factors** that provide an easy and compact form of handling **categorical (nominal)** data. Factors have **levels** and particularly useful in datasets where you have nominal variables with a fixed number of possible values. R stores these values **as numeric codes** that are considerably more memory efficient.

Suppose you have a vector with the gender of ten individuals. We can transform this vector into a factor:

```
>g <- c("f", "m", "m", "m", "f", "m", "f", "m", "f", "f")
> g <- factor(g)
>g
```

Output:

```
[1] f m m m f m f m f f
```

Levels: f m

To find out **what data types** are, "**class()**" function is used **in R** and "**type()**" function is used **in Python**.

```
> print(class(a)) # R
[1] "numeric"
print(type(a)) # Python
<type 'list'>
```

## Part 3 - Reading Data

**Public datasets** are available in the **R base** and **Python Scikit-Learn** library. You can check built-in datasets in R with "**data()**" command. We will use the **iris flower dataset** as an example. It is often used for testing out machine learning algorithms and visualizations. This dataset contains 3 classes of 50 instances each, where each class refers to a type of iris plant. You can show the iris data by typing "**iris**" keyword in R.

In Python:

```
import pandas as pd
from sklearn import datasets

# This creates a Sklearn bunch
data = datasets.load_iris()
# Convert to Pandas dataframe
iris = pd.DataFrame(data.data, columns=data.feature_names)
```

In R:

```
# it can be obtained only typing iris
iris
```

You can obtain information about data, show row and column (feature) sizes with **str()** and **dim()** in R, **info()** and **shape** in Python. First 3 and last 3 rows are listed with **head(3)** and **tail(3)** commands.

In Python:

```
print(iris.info())
print(iris.shape)
print(iris.head(3))
print(iris.tail(3))
```

In R:

```
str(iris)
dim(iris)
head(df, 3)
tail(df, 3)
```

Output:

```
> str(iris)
'data.frame': 150 obs. of 5 variables:
 $ Sepal.Length: num  5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
 $ Sepal.Width : num  3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
 $ Petal.Length: num  1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
 $ Petal.Width : num  0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
 $ Species     : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1 1 1 1 1 1 1 ...

> dim(iris)
[1] 150  5

> head(df, 3)
  Transaction_date Product Price Payment_Type      Name
1   1/2/09 6:17 Product1 1200  Mastercard    carolina
2   1/2/09 4:53 Product1 1200      Visa    Betina Parkville
3   1/2/09 13:08 Product1 1200  Mastercard Federica e Andrea Astoria

> tail(df, 3)
  Transaction_date Product Price Payment_Type      Name
996   1/1/09 4:24 Product3 7500      Amex Pamela Skaneateles
997   1/8/09 11:55 Product1 1200      Diners julie
998   1/12/09 21:30 Product1 1200      Visa Julia Madison
```

If we want to get some **subsets** in the data, we can use the **indexes** or **subset** function in R. If we want to obtain data with the number of sepal length is greater than 5, the code can be written as follows.

In Python:

```
print(len(iris[iris['sepal length (cm)']>5.0]))
```

In R:

```
nrow(iris[iris$Sepal.Length>5.0,])  
  
# with subset  
nrow(subset(iris, iris$Sepal.Length>5.0))
```

Output:

118

**Total** and **average** values can be calculated with the code below.

In Python:

```
def calc_mean(col_name):  
    result = iris[col_name].mean()  
  
    return result  
  
def calc_mean_2(col_name):  
    result = iris[col_name].sum() / len(iris[col_name])  
  
    return result  
  
print(calc_mean("sepal width (cm)"))  
print(calc_mean_2("sepal width (cm)"))
```

In R:

```
calc_mean <- function(col_name){  
    result <- mean(col_name)  
  
    return(result)  
}
```

```
calc_mean_2 <- function(col_name){  
  result <- sum(col_name) / length(col_name)  
  return(result)  
}  
  
calc_mean(iris$Sepal.Width)  
calc_mean_2(iris$Sepal.Width)
```

Output:

```
3.057333  
3.057333
```

**Comma-separated values (CSV)** data is widely used in data mining. **Download** the Amazon sales **CSV data** which was prepared for this course, called **AmazonSales.csv**. A command like the following can be written in both R and Python to read the data. Data will be obtained in **dataframe** structure.

**Exercise 1.** Read the data separately in Python and R.

In Python:

In R:

**Exercise 2.** How many countries are there in the data? Write the output and commands separately in Python and R.

In Python:

In R:

Output:

**Exercise 3.** Find the total sales amount for each product. Write the output and commands separately in Python and R.

In Python:

In R:

Output:

**Exercise 4.** Find the total fee charged with Mastercard. Write the output and commands separately in Python and R.

In Python:

In R:

Output:

## Part 4 - Visualization Data

Visualization may be required to understand the data. In this way, information is obtained about data. **Matplotlib** library is widely used to visualize data with **Python** whereas **base** plotting functions can be used for visualizing with **R**. **Box plot** is a type of graph that shows how the values in the data are **spread out**. Data is expressed in **quarters** as minimum, first quartile (25%), median, third quartile (50%), and maximum values.

In Python:

```
import matplotlib.pyplot as plt
```


```
img=plt.boxplot(iris['sepal length (cm)'])  
plt.show(img)
```



In R:

```
boxplot(iris$Sepal.Length)
```

Output:



**Scatter plot** is used to determine the relationship between two different variables. It can be seen whether there is a direct **relationship between** the **two variables** and how strong this relationship is. In addition, **changes in time series** data can be observed by "**type**" information in R and "**plot()**" function in Python.

In Python:

```
import matplotlib.pyplot as plt
```

```
img=plt.scatter(iris["sepal width (cm)"], iris["sepal length (cm)"])  
plt.show(img)
```

```
# Time series plot
```

```
img=plt.plot(iris["sepal length (cm)"])  
plt.show(img)
```

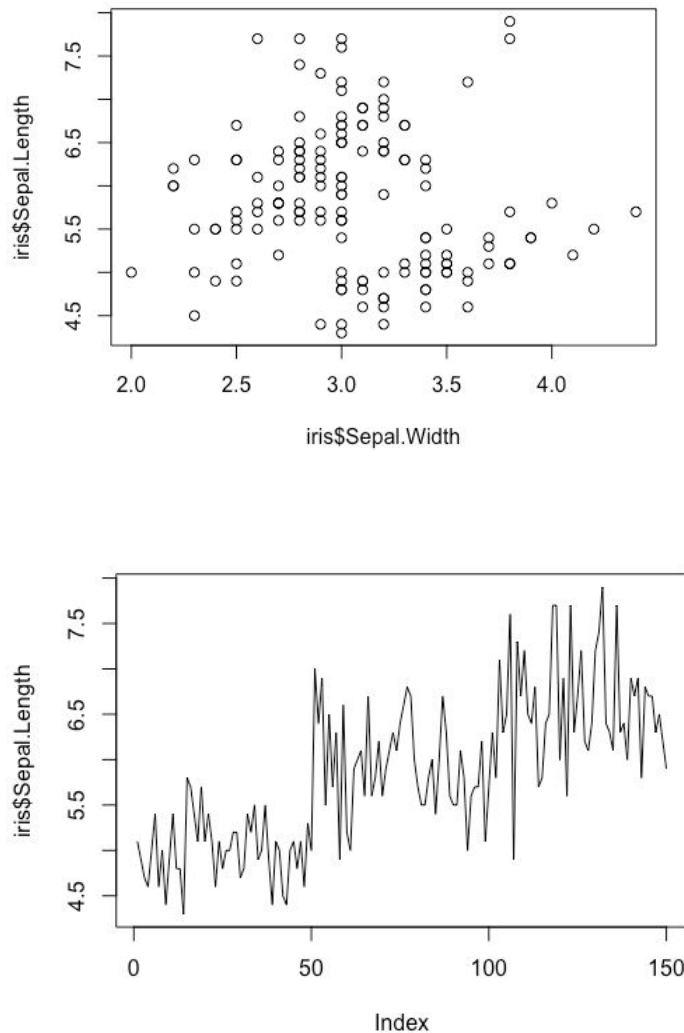
In R:

```
plot(iris$Sepal.Width,iris$Sepal.Length)
```

```
# Time series plot
```

```
plot(iris$Sepal.Length, type = 'l')
```

Output:



**Histogram plot** is used to determine the **distribution** and **frequency** of the data. It divides values into **groups** and returns the **frequency**.


In Python:

```
import matplotlib.pyplot as plt  
  
img=plt.hist(iris["sepal length (cm)"])  
plt.title('Sepal Histogram')  
plt.show()
```

In R:

```
hist(iris$Sepal.Length, col="aliceblue", main="Sepal Histogram")
```

Output:



**Exercise 5.** Draw scatter prices by product and compare prices from Amazon Sales CSV file. Write the output and commands separately in Python and R.

In Python:

In R:

Output:

**Exercise 6.** Describe how you would create visualizations to display information that describes the following types of systems.

- (a) Computer **networks**. Be sure to include both the static aspects of the network, such as connectivity, and the dynamic aspects, such as traffic.
- (b) The **distribution** of specific plant and animal species around the world for a specific moment in time.
- (c) The use of computer **resources**, such as processor time, main memory, and disk, for a set of benchmark database programs.
- (d) The change in **occupation** of workers in a particular country over the last thirty years. Assume that you have yearly information about each person that also includes gender and level of education.

Answer:

- a)
- b)
- c)
- d)

## Part 5 - Applying Data Mining

In this section, we will write Python and R code to **classify** species in the iris dataset. In the reading data section, the subject of how the data will be obtained is mentioned. The data obtained should be prepared for the application of mathematical methods. **Features** and **classes** should be determined before performing classification. Iris dataset consists of **3 labels**

Page 10 of 10

## In Python:

```
from sklearn import datasets
```

```
# Load built-in dataset
iris = datasets.load_iris()
```

```
x = iris.data
y = iris.target
y_labels = iris.target_names
```

```
print(y)
print(y_labels)
```

In R:

```
data(iris)

x <- iris[1:4]
y <- iris$Species
y_labels <- levels(iris$Species)

print(y)
print(y_labels)
```

Output:

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1  
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2 2 2  
2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2  
2 2]  
['setosa' 'versicolor' 'virginica']
```

**Standardization** or **normalization** will be applied to make an entire set of values have a particular property. If **different variables** are to be combined in some way, then such a transformation is often necessary to avoid having a variable with **large values dominate** the results of the calculation. For example; The **Gaussian normalization** ( $val = (x - \text{mean}(x)) / \text{std}(x)$ ) creates a new variable that has a mean of 0 and a standard deviation of 1. **Min-max** normalization is applied in the codes below.

In Python:

```
from sklearn.preprocessing import MinMaxScaler
import matplotlib.pyplot as plt

min_max_scaler = MinMaxScaler()
x_norm = min_max_scaler.fit_transform(x)

img=plt.boxplot(x)
plt.show()
img=plt.boxplot(x_norm)
plt.show()
```


In R:

```
# Build your min-max function
normalize <- function(x) {
  num <- x - min(x)
  denom <- max(x) - min(x)
  return (num/denom)
}

# Normalize the data
x_norm <- as.data.frame(lapply(x, normalize))

# Check the normalization
boxplot(x)
boxplot(x_norm)
```

Output:



Data will be divided into training and test sets to perform training and **evaluate the performance** of the models developed.

In Python:

```
from sklearn.model_selection import train_test_split

# set the random state to make reproducible
x_train,x_test,y_train,y_test=train_test_split(x, y, test_size=.25, random_state=123)

# Check data size
print(x_train.shape)
print(x_test.shape)
```

In R:

```
# set the seed to make reproducible
set.seed(123)

train_ind <- sample(seq_len(nrow(x_norm)), size = floor(0.75 * nrow(x_norm)))

x_train = x_norm[train_ind,]
x_test = x_norm[-train_ind,] # all indices except train indexes
y_train = y[train_ind]
y_test = y[-train_ind]

# Check data size
dim(x_train)
dim(x_test)
```

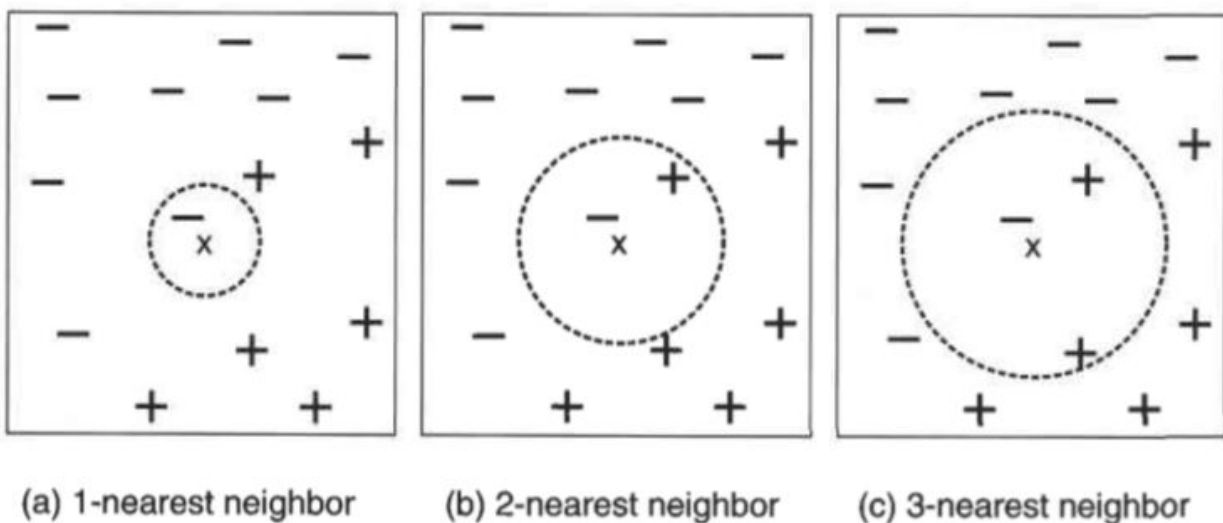


Output:

(112, 4)

(38, 4)

**KNN (K-Nearest-Neighbor)** model will be implemented in two language. The algorithm computes the **distance** (or similarity) between each test example and all the training examples to determine its nearest-neighbor list. We will use the **knn()** function in **R** and **KNeighborsClassifier()** function in **Python**, both of them use the **Euclidian distance** measure in order to find the k-nearest neighbours to your new, unknown instance. Here, the **k parameter** is one that you set yourself.



In Python:

```
from sklearn import neighbors
from sklearn.metrics import accuracy_score, confusion_matrix

classifier=neighbors.KNeighborsClassifier(n_neighbors=5)

classifier.fit(x_train,y_train)

iris_pred = classifier.predict(x_test)

print(accuracy_score(y_test, iris_pred))
```

```
print(confusion_matrix(y_test, iris_pred))
```

In R:

```
install.packages("Metrics") # For accuracy
install.packages("caret") # For confusion matrix
install.packages('e1071', dependencies=TRUE) # For caret

library(class) # For knn()
library(Metrics)
library(caret)

iris_pred <- knn(train = x_train, test = x_test, cl = y_train, k=5)
iris_pred

accuracy(iris_pred, y_test)

# Confusion matrix
table(iris_pred, y_test)

# Confusion matrix with overall statistics
confusionMatrix(iris_pred, y_test)
```

Output:

```
0.9736842105263158
```

```
[[16 0 0]
 [ 0 7 1]
 [ 0 0 14]]
```

## Part 6 - Assignment

In this assignment, the **following tasks** will be done:

- Write the most appropriate answers to the **exercises (Exercise 1, 2, 3, 4, 5, 6)** in the document.
- In addition, you will download the **breast cancer dataset** and write the Python and R code that can analyse the data.
- Which **visualization** techniques would be more appropriate for analyzing this data (breast cancer)? Apply many different visualization techniques.
- What should we do to **find outliers** in the data? Explain and apply your solution.
- You should **decide** whether **pre-processing** is necessary for this data. Explain which technique is appropriate for the data.
- What **type of data mining** (classification, clustering, etc.) you think would be relevant? Apply the data mining technique that you decide is relevant.

### Notes:

- You can download the dataset as a CSV file from **UCI** (University of California, Irvine), “wdbc.data” will be used for analysing (<https://archive.ics.uci.edu/ml/machine-learning-databases/breast-cancer-wisconsin/wdbc.data>).
- Please upload your **Python, R codes** and **your report** (with answers to questions from **exercises and tasks**) to CMS until **20 March 2020 23:00**. You should upload a zip file, and file name “Student\_Number\_Name.zip”.

---

## References

- R and Data Mining: Examples and Case Studies, Yanchang Zhao, 2015.
- Data Mining with R, Learning with Case Studies, Chapman & Hall/CRC Data Mining and Knowledge Discovery Series SERIES EDITOR, Vipin Kumar.
- Learning Data Mining with Python, Second Edition, Robert Layton Book.
- <http://www.cs.ukzn.ac.za/~hughm/dm/content/slides01.pdf>
- <https://cran.r-project.org/doc/manuals/r-release/R-intro.pdf>
- [https://shiring.github.io/r\\_vs\\_python/2017/01/22/R\\_vs\\_Py\\_post](https://shiring.github.io/r_vs_python/2017/01/22/R_vs_Py_post)
- <https://www.datacamp.com/community/tutorials/machine-learning-in-r>
- <https://www.analyticsvidhya.com/blog/2016/02/complete-tutorial-learn-data-science-scratch/>
- <https://docs.python.org/3/library/datetime.html>
- [https://www.tutorialspoint.com/python3/python\\_tutorial.pdf](https://www.tutorialspoint.com/python3/python_tutorial.pdf)
- <https://www.jetbrains.com/help/pycharm/installation-guide.html>
- <https://beginnersbook.com/>